

# Text Categorization: Discriminative Classifiers

## Part 1

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# Overview

- What is text categorization?
- Why text categorization?
- How to do text categorization?
  - Generative probabilistic models
  - **Discriminative approaches**
- How to evaluate categorization results?

# Anatomy of Naïve Bayes Classifier

Two categories:  $\theta_1$  and  $\theta_2$

$$\text{score}(d) = \log \frac{p(\theta_1 | d)}{p(\theta_2 | d)} = \log \frac{p(\theta_1) \prod_{w \in V} p(w | \theta_1)^{c(w,d)}}{p(\theta_2) \prod_{w \in V} p(w | \theta_2)^{c(w,d)}}$$

$$= \log \frac{p(\theta_1)}{p(\theta_2)} + \sum_{w \in V} c(w,d) \log \frac{p(w | \theta_1)}{p(w | \theta_2)}$$

Category bias ( $\beta_0$ ) doesn't depend on  $d$ !

Sum over all words (features  $\{x_i\}$ )

Weight on each word (feature)  $\beta_i$

Feature value:  $x_i = c(w,d)$



Generalize

$$d = (x_1, x_2, \dots, x_M), \quad x_i \in \mathcal{R}$$

$$\text{score}(d) = \beta_0 + \sum_{i=1}^M x_i \beta_i \quad \beta_i \in \mathcal{R}$$

= Logistic Regression!

# Discriminative Classifier 1: Logistic Regression

**Binary Response Variable:**  $Y \in \{0,1\}$

**Predictors:**  $X = (x_1, x_2, \dots, x_M)$ ,  $x_i \in \mathbb{R}$

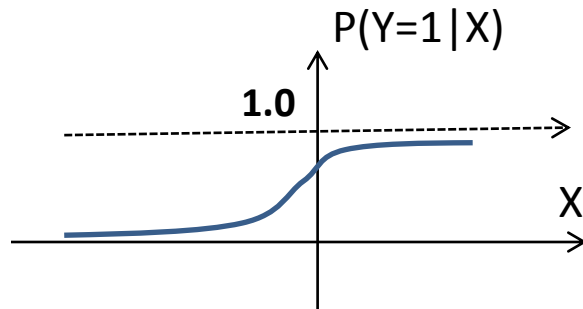
$$Y = \begin{cases} 1 & \text{category}(d) = \theta_1 \\ 0 & \text{category}(d) = \theta_2 \end{cases}$$

**Modeling  $p(Y|X)$  directly**

**Allow many other features than words!**

$$\log \frac{p(\theta_1 | d)}{p(\theta_2 | d)} = \log \frac{p(Y = 1 | X)}{p(Y = 0 | X)} = \log \frac{p(Y = 1 | X)}{1 - p(Y = 1 | X)} = \beta_0 + \sum_{i=1}^M x_i \beta_i \quad \beta_i \in \mathbb{R}$$

$$p(Y = 1 | X) = \frac{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i}}{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i} + 1}$$



# Estimation of Parameters

- Training Data:  $T = \{(X_i, Y_i)\}, i=1, 2, \dots, |T|$
- Parameters:  $\vec{\beta} = (\beta_0, \beta_1, \dots, \beta_M)$
- Conditional likelihood:  $p(T | \vec{\beta}) = \prod_{i=1}^{|T|} p(Y = Y_i | X = X_i, \vec{\beta})$

$Y_i = 1$

$$p(Y = 1 | X) = \frac{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i}}{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i} + 1}$$

$Y_i = 0$

$$p(Y = 0 | X) = \frac{1}{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i} + 1}$$

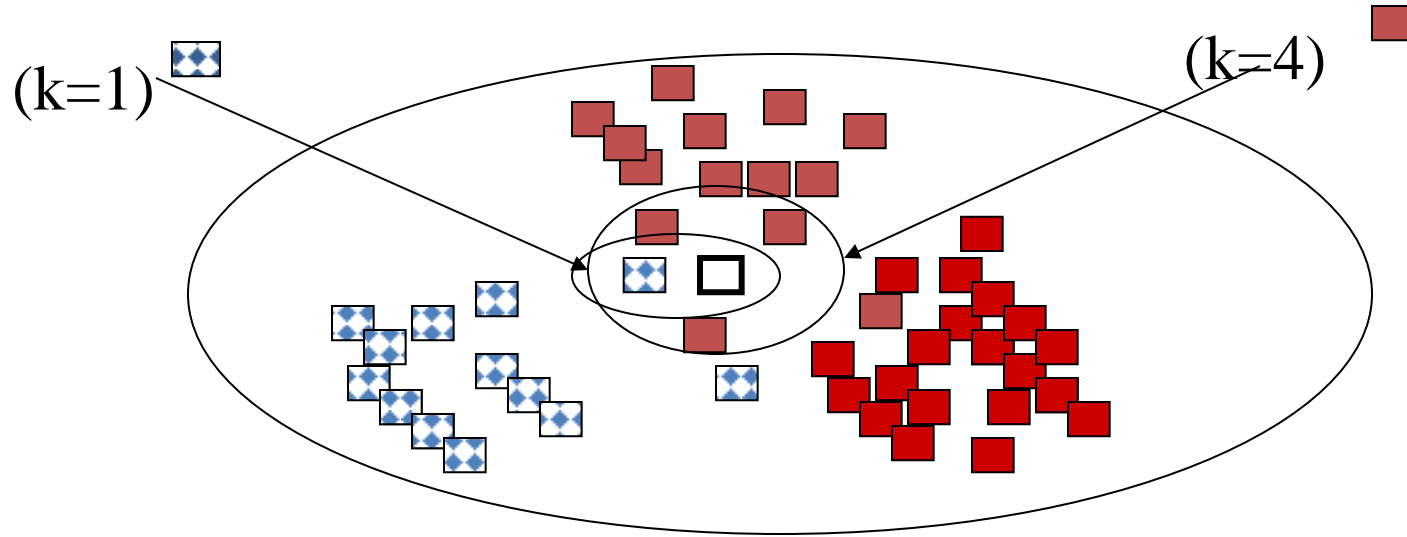
- Maximum Likelihood estimate  $\vec{\beta}^* = \arg \max_{\vec{\beta}} p(T | \vec{\beta})$

Can be computed in many ways (e.g., Newton's method)

## Discriminative Classifier 2: K-Nearest Neighbors (K-NN)

- Find  $k$  examples in the training set that are most similar to the text object to be classified (“neighbor” documents)
- Assign the category that is most common in these neighbor text objects (neighbors vote for the category)
- Can be improved by considering the distance of a neighbor (a closer neighbor has more influence)
- Can be regarded as a way to directly estimate the conditional probability of label given data instance, i.e.,  $p(Y|X)$
- Need a similarity function to measure similarity of two text objects

# Illustration of K-NN Classifier




# K-NN as an Estimate of $p(Y|X)$

Assume  $p(\theta_i | d)$  is locally smooth, i.e.,  
the same for all the  $d$ 's in this region  $R$



$$p(\theta_i | d) = p(\theta_i | R)$$



Estimate  $p(\theta_i | R)$  based on  
the known categories in the region

$$p(\theta_i | R) = \frac{c(\theta_i, R)}{|R|}$$

Count of  $d$ 's in  $R$   
with category  $\theta_i$

Total # of  
docs in  $R$



# Text Categorization: Discriminative Classifiers

## Part 2

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

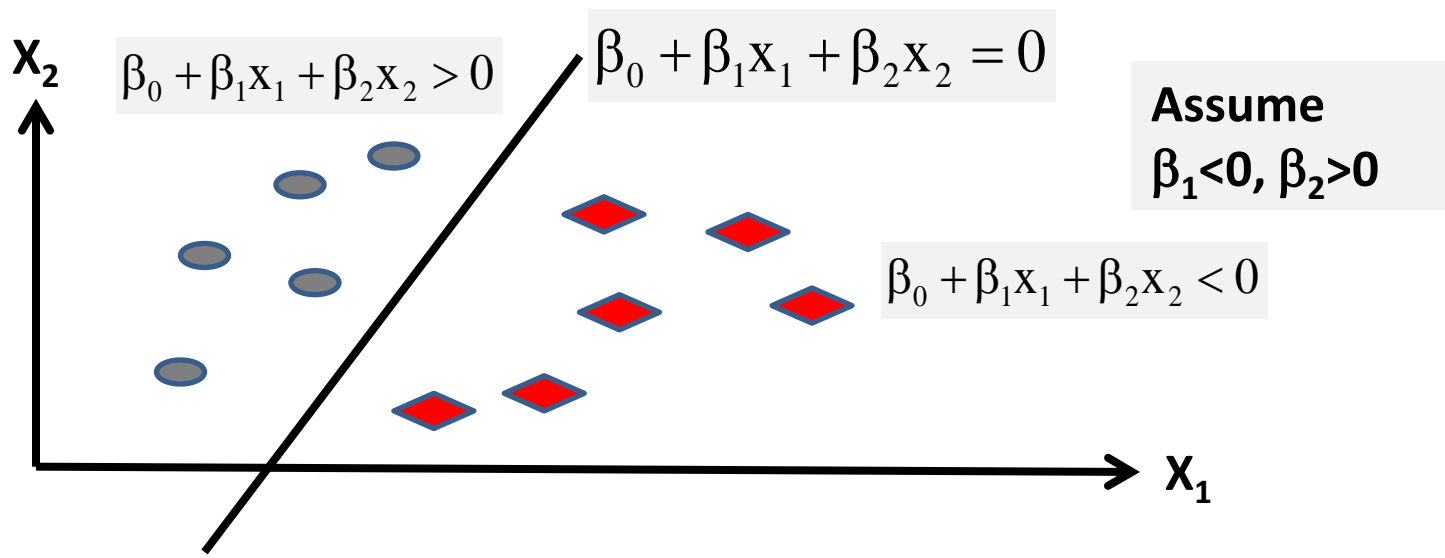
# Discriminative Classifier 3: Support Vector Machine (SVM)

- Consider two categories:  $\{\theta_1, \theta_2\}$

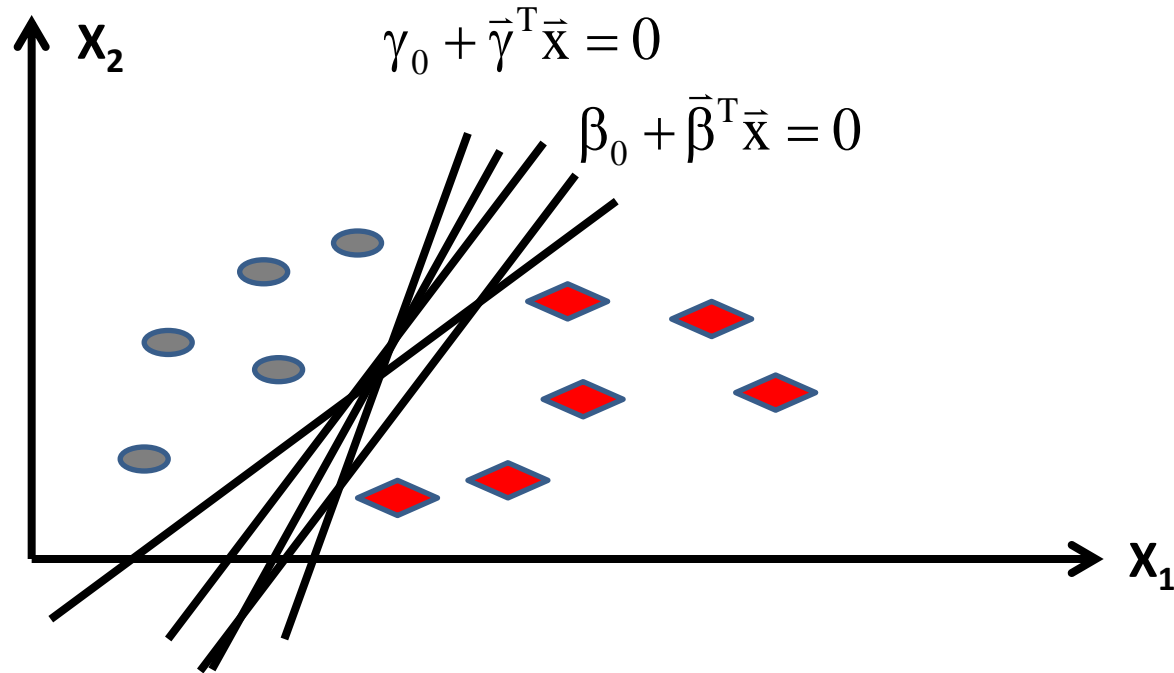
$f(X) \geq 0 \Rightarrow X$  is in category  $\theta_1$

$f(X) < 0 \Rightarrow X$  is in category  $\theta_2$

- Use a linear separator  $f(X) = \beta_0 + \sum_{i=1}^M x_i \beta_i \quad \beta_i \in \mathbb{R}$



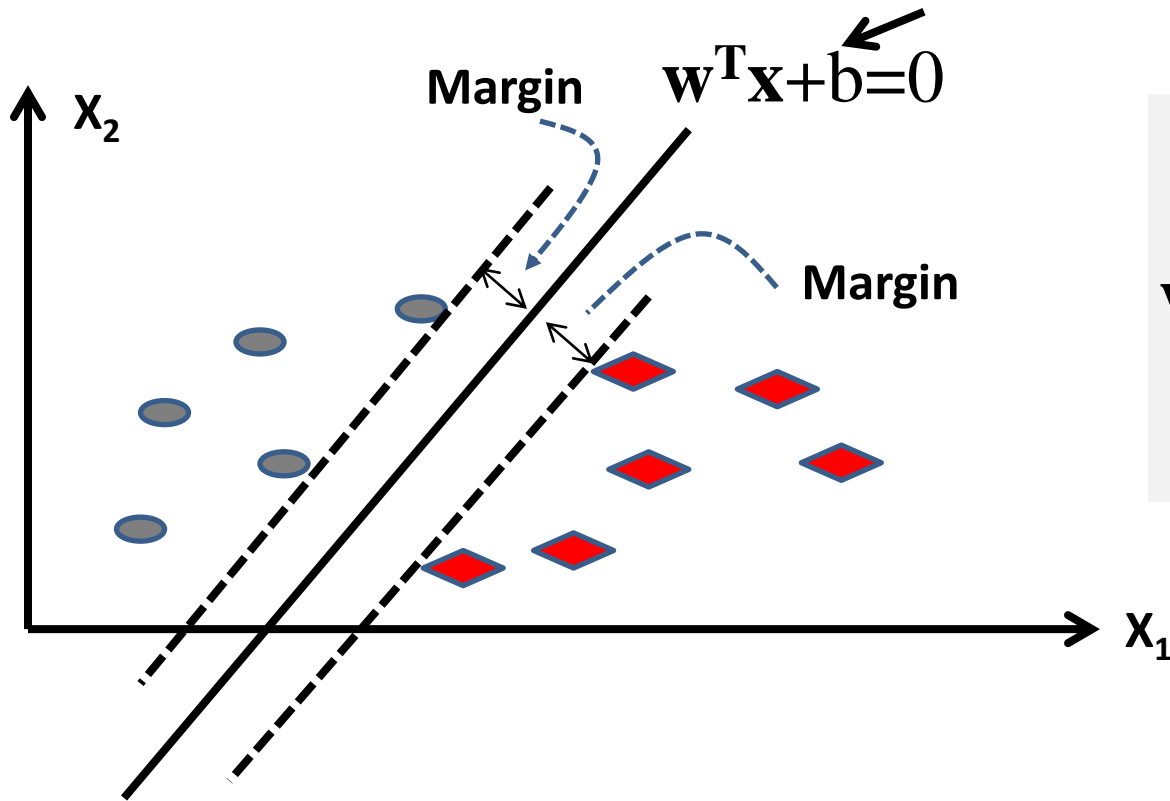
# Which Linear Separator Is the Best?



# Best Separator = Maximize the Margin

Notation Change:  $\beta \rightarrow w$ ;  $\beta_0 \rightarrow b$

Bias constant



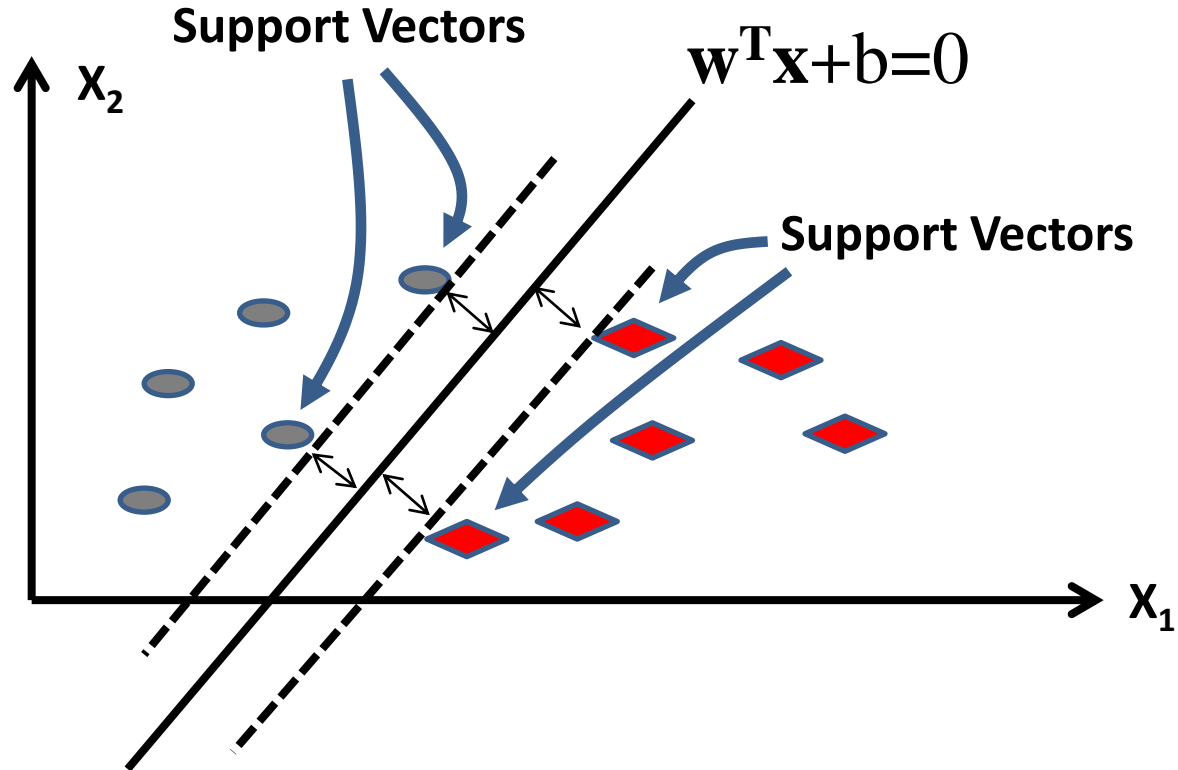
Feature Weights

$$\mathbf{w} = \begin{pmatrix} w_1 \\ w_2 \\ \dots \\ w_M \end{pmatrix}$$

Feature Vector  
(e.g., word counts)

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \dots \\ x_M \end{pmatrix}$$

# Only the Support Vectors Matter



# Linear SVM

**Classifier:**  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$

**Parameters:**  $\mathbf{w}$ ,  $b$

**Training Data:**  $T = \{(\mathbf{x}_i, y_i)\}$ ,  $i=1, \dots, |T|$ .  $\mathbf{x}_i$  is a feature vector;  $y_i \in \{-1, 1\}$

$f(X) \geq 0 \Rightarrow X$  is in category  $\theta_1$

$f(X) < 0 \Rightarrow X$  is in category  $\theta_2$

**Goal 1: Correct labeling on training data:**

If  $y_i = 1 \rightarrow \mathbf{w}^T \mathbf{x}_i + b \geq 1$

If  $y_i = -1 \rightarrow \mathbf{w}^T \mathbf{x}_i + b \leq -1$



Constraint

$$\forall i, y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

Objective

$$\text{Minimize } \Phi(\mathbf{w}) = \mathbf{w}^T \mathbf{w}$$

**Goal 2: Maximize margin**

Large margin  $\Leftrightarrow$  Small  $\mathbf{w}^T \mathbf{w}$

The optimization problem is quadratic programming with linear constraints

# Linear SVM with Soft Margin

**Classifier:**  $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b > 0$ ?

**Parameters:**  $\mathbf{w}$ ,  $b$

**Training Data:**  $T = \{(\mathbf{x}_i, y_i)\}, i=1, \dots, |T|$ .

**Find  $\mathbf{w}$ ,  $b$ , and  $\xi_i$  to minimize**  $\Phi(\mathbf{w}) = \mathbf{w}^T \mathbf{w} + C \sum_{i \in [1, |T|]} \xi_i$

Added to allow training errors

**Subject to**  $\forall i \in [1, |T|], y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \xi_i \geq 0$

$C > 0$  is a parameter to control the trade-off between minimizing the errors and maximizing the margin

**The optimization problem is still quadratic programming with linear constraints**

# Summary of Text Categorization Methods

- Many methods are available, but no clear winner
  - All require effective feature representation (need domain knowledge)
  - It is useful to compare/combine multiple methods for a particular problem
- Most techniques rely on supervised machine learning and thus can be applied to **any** text categorization problem!
  - Humans annotate training data and design features
  - Computer optimizes the combination of features
  - Good performance requires 1) effective features and 2) plenty of training data
  - Performance is generally (much) more affected by the effectiveness of features than by the choice of a specific classifier



# Summary of Text Categorization Methods (cont.)

- How to design effective features? (application-specific)
  - Analyze the categorization problem and exploit domain knowledge
  - Perform error analysis to obtain insights
  - Leverage machine learning techniques (e.g., feature selection, dimension reduction, deep learning)
- How to obtain “enough” training examples?
  - Low-quality (“pseudo”) training examples may be leveraged
  - Exploit unlabeled data (using semi-supervised learning techniques)
  - Domain adaptation/transfer learning (“borrow” training examples from a related domain/problem)

# Suggested Reading

Manning, Chris D., Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. Cambridge: Cambridge University Press, 2007.  
(Chapters 13-15)

# Text Categorization: Evaluation

## Part 1

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# Overview

- What is text categorization?
- Why text categorization?
- How to do text categorization?
  - Generative probabilistic models
  - Discriminative approaches
- **How to evaluate categorization results?**

# General Evaluation Methodology

- Have humans to create a test collection where every document is tagged with the desired categories (“ground truth”)
- Generate categorization results using a system on the test collection
- Compare the system categorization decisions with the human-made categorization decisions and quantify their similarity (or equivalently difference)
  - The higher the similarity is, the better the results are
  - Similarity can be measured from different perspectives to understand the quality of results in detail (e.g., which category performs better?)
  - In general, different categorization mistakes may have a different cost that inevitably depends on specific applications, but it is okay not to consider such a cost variation for **relative comparison of methods**

# Classification Accuracy (Percentage of Correct Decisions)

	$\mathbf{c_1}$	$\mathbf{c_2}$	$\mathbf{c_3}$	$\mathbf{\dots}$	$\mathbf{c_k}$	
$\mathbf{d_1}$	y(+)	y(-)	n(+)		n(+)	+/- <b>human answer</b>
$\mathbf{d_2}$	y(-)	n(+)	y(+)		n(+)	(+= correct; - =incorrect)
$\mathbf{d_3}$	n(+)	n(+)	y(+)		n(+)	<b>y/n system result</b>
$\mathbf{\dots}$						(y=yes; n=no)
$\mathbf{d_N}$	$\mathbf{\dots}$	$\mathbf{\dots}$				

$$\begin{aligned}
 \text{Classification Accuracy} &= \frac{\text{Total number of correct decisions}}{\text{Total number of decisions made}} \\
 &= \frac{\text{count}(y(+)) + \text{count}(n(-))}{kN}
 \end{aligned}$$

# Problems with Classification Accuracy

- Some decision errors are more serious than others
  - It may be more important to get the decisions right on some documents than others
  - It may be more important to get the decisions right on some categories than others
  - E.g., spam filtering: missing a legitimate email costs more than letting a spam go
- Problem with imbalanced test set
  - Skewed test set: 98% in category 1; 2% in category 2
  - Strong baseline: put all instances in category 1 → 98% accuracy!

# Per-Document Evaluation

	$c_1$	$c_2$	$c_3$	$\dots$	$c_k$
$d_1$	y(+)	y(-)	n(+)		n(+)
$d_2$	y(-)	n(+)	y(+)		n(+)
$d_3$	n(+)	n(+)	y(+)		n(+)

How good are the decisions on  $d_i$ ?

When the system says “yes,”  
how many are correct?

↓

Precision =  $\frac{TP}{TP + FP}$

↗

Recall =  $\frac{TP}{TP + FN}$

	System (“y”)	System (“n”)
Human (+)	True Positives TP	False Negatives FN
Human (-)	False Positives FP	True Negatives TN

Does the doc have all the categories  
it should have?



# Per-Category Evaluation

	$c_1$	$c_2$	$c_3$	...	$c_k$
$d_1$	y(+)	y(-)	n(+)		n(+)
$d_2$	y(-)	n(+)	y(+)		n(+)
$d_3$	n(+)	n(+)	y(+)		n(-)

How good are the decisions on  $c_i$ ?

When the system says “yes,”  
how many are correct?

↓

**Precision** =  $\frac{TP}{TP + FP}$

↗

**Recall** =  $\frac{TP}{TP + FN}$

	System (“y”)	System (“n”)
Human (+)	True Positives TP	False Negatives FN
Human (-)	False Positives FP	True Negatives TN

Has the category been assigned to  
all the docs of this category?

# Combine Precision and Recall: F-Measure

$$F_{\beta} = \frac{1}{\frac{\beta^2}{\beta^2+1} \frac{1}{R} + \frac{1}{\beta^2+1} \frac{1}{P}} = \frac{(\beta^2 + 1)P * R}{\beta^2 P + R}$$

$$F_1 = \frac{2PR}{P + R}$$

**P**: precision

**R**: recall

**β**: parameter (often set to 1)

Why not  $0.5 * P + 0.5 * R$ ?

What is R if the system says “y” for all category-doc pairs?

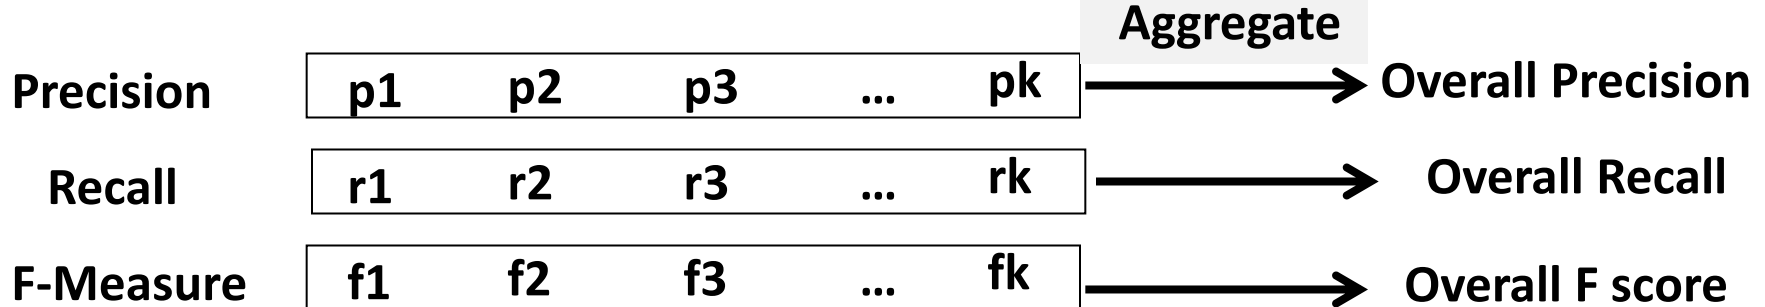
# Text Categorization: Evaluation

## Part 2

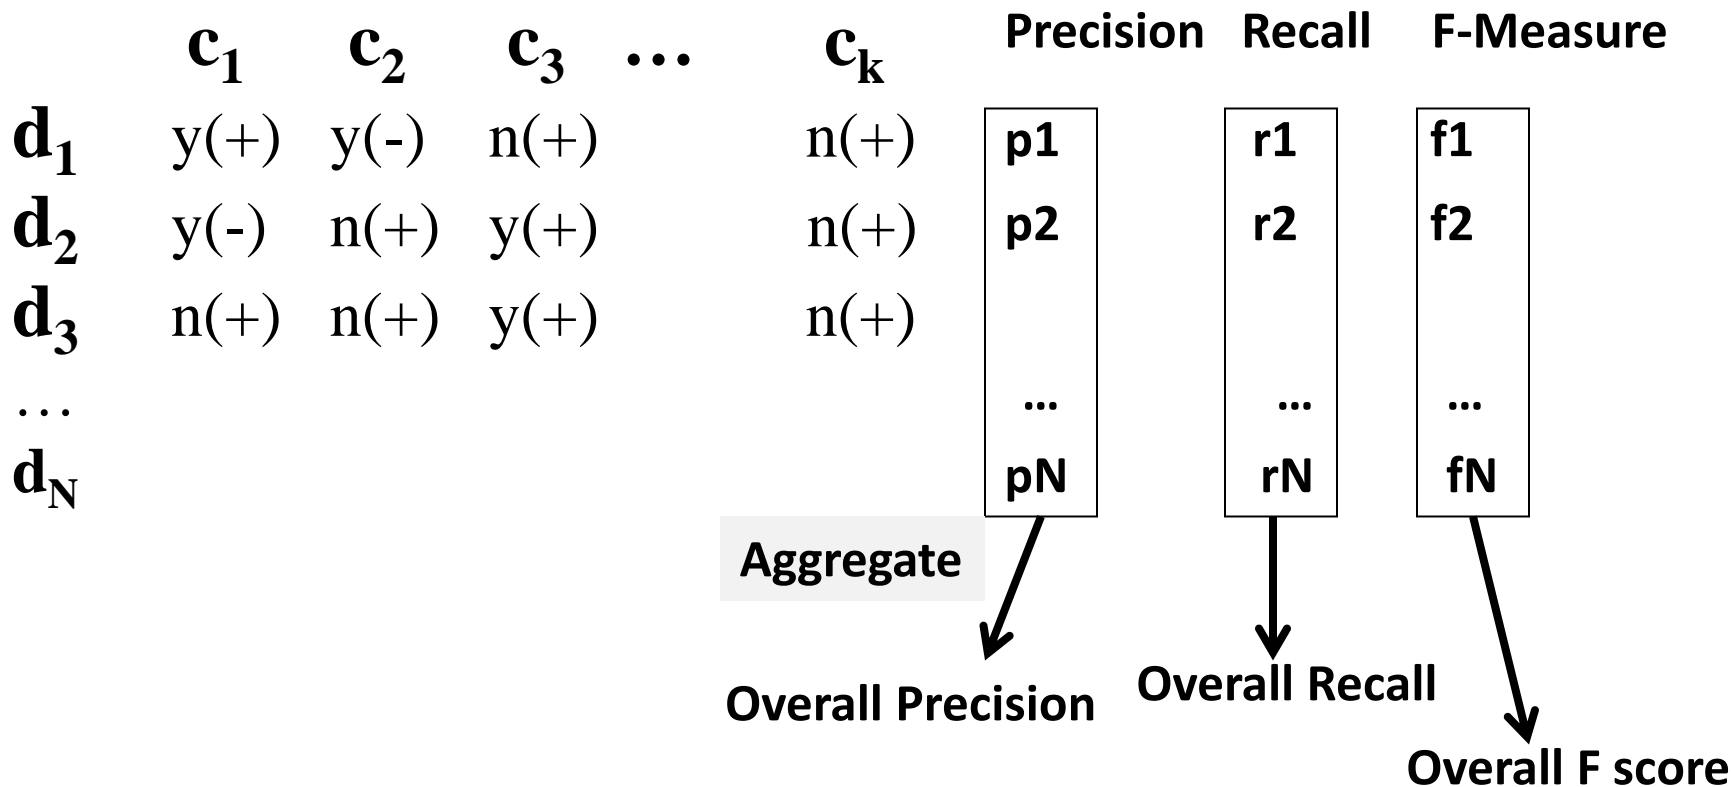
ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# (Macro) Average Over All the Categories

	$c_1$	$c_2$	$c_3$	...	$c_k$
$d_1$	y(+)	y(-)	n(+)		n(+)
$d_2$	y(-)	n(+)	y(+)		n(+)
$d_3$	n(+)	n(+)	y(+)		n(+)
...					
$d_N$	...	...			



# (Macro) Average Over All the Documents



# Micro-Averaging of Precision and Recall

	$c_1$	$c_2$	$c_3$	$\dots$	$c_k$
$d_1$	y(+)	y(-)	n(+)		n(+)
$d_2$	y(-)	n(+)	y(+)		n(+)
$d_3$	n(+)	n(+)	y(+)		n(+)
$\dots$					
$d_N$	...	...			

First pool all decisions,  
then compute precision and recall



$$\text{Precision} = \frac{TP}{TP + FP}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

	System ("y")	System ("n")
Human (+)	True Positives ( TP)	False Negatives (FN)
Human (-)	False Positives(FP)	True Negatives(TN)

# Sometimes Ranking Is More Appropriate

- The categorization results are often passed to a human for
  - further editing (e.g., correcting system mistakes on news categories)
  - prioritizing a task (e.g., routing an email to the right person for processing)
- In such cases, we can evaluate the results as a ranked list if the system can give scores for the decisions
  - E.g., discovery of spam emails (➔ rank emails for the “spam” category)
  - Often more appropriate to frame the problem as a ranking problem instead of a categorization problem (e.g., ranking documents in a search engine)

# Summary of Categorization Evaluation

- Evaluation is always very important, so get it right!
- Measures must reflect the **intended use** of the results for a particular application (e.g., spam filtering vs. news categorization)
  - Consider: How will the results be further processed (by a user)?
  - Ideally associate a different cost with each different decision error
- Commonly used measures for **relative** comparison of different methods:
  - Accuracy, precision, recall, F score
  - Variations: per-document, per-category, micro vs. macro averaging
- Sometimes **ranking** may be more appropriate



# Suggested Reading

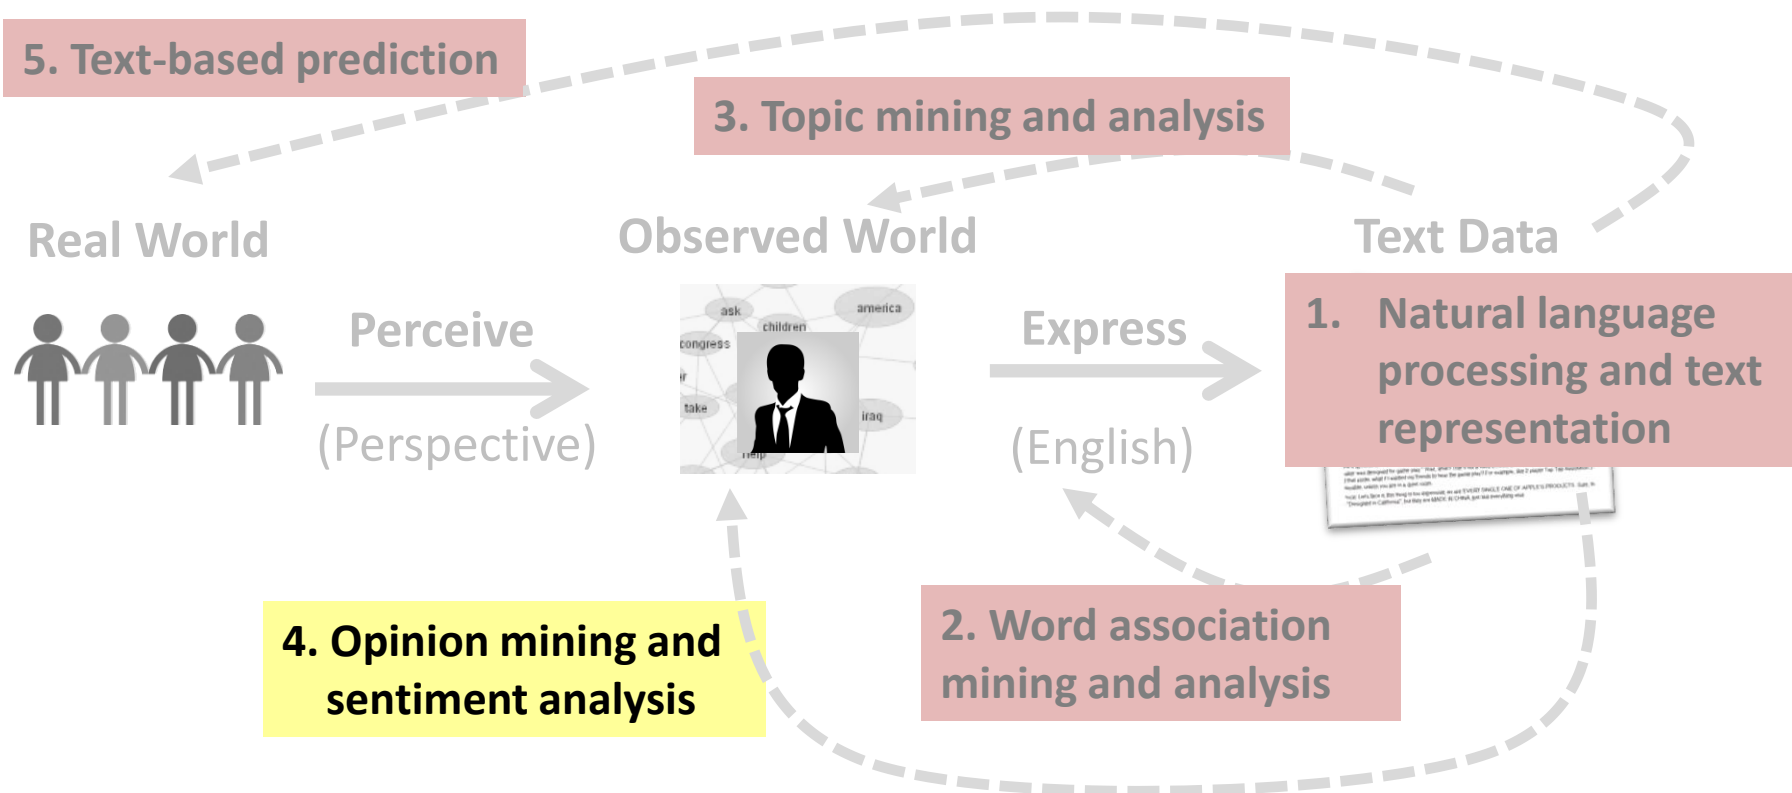
- Manning, Chris D., Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. Cambridge: Cambridge University Press, 2007. (Chapters 13-15)
- Yang, Yiming. 1999. An Evaluation of Statistical Approaches to Text Categorization. *Inf. Retr.* 1, 1-2 (May 1999), 69-90. DOI=10.1023/A:1009982220290



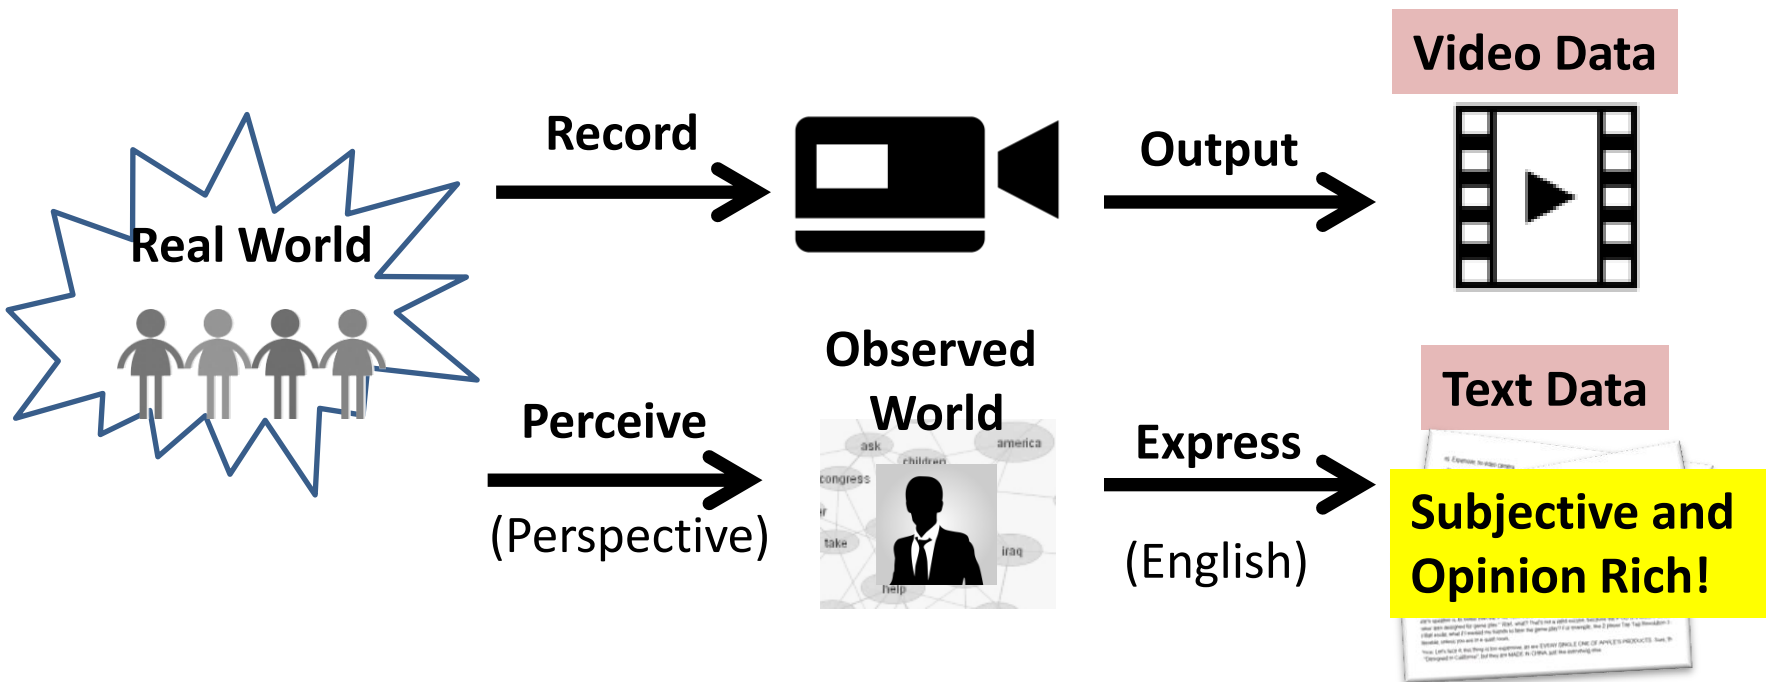
# Opinion Mining and Sentiment Analysis: Motivation

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# Opinion Mining and Sentiment Analysis: Motivation



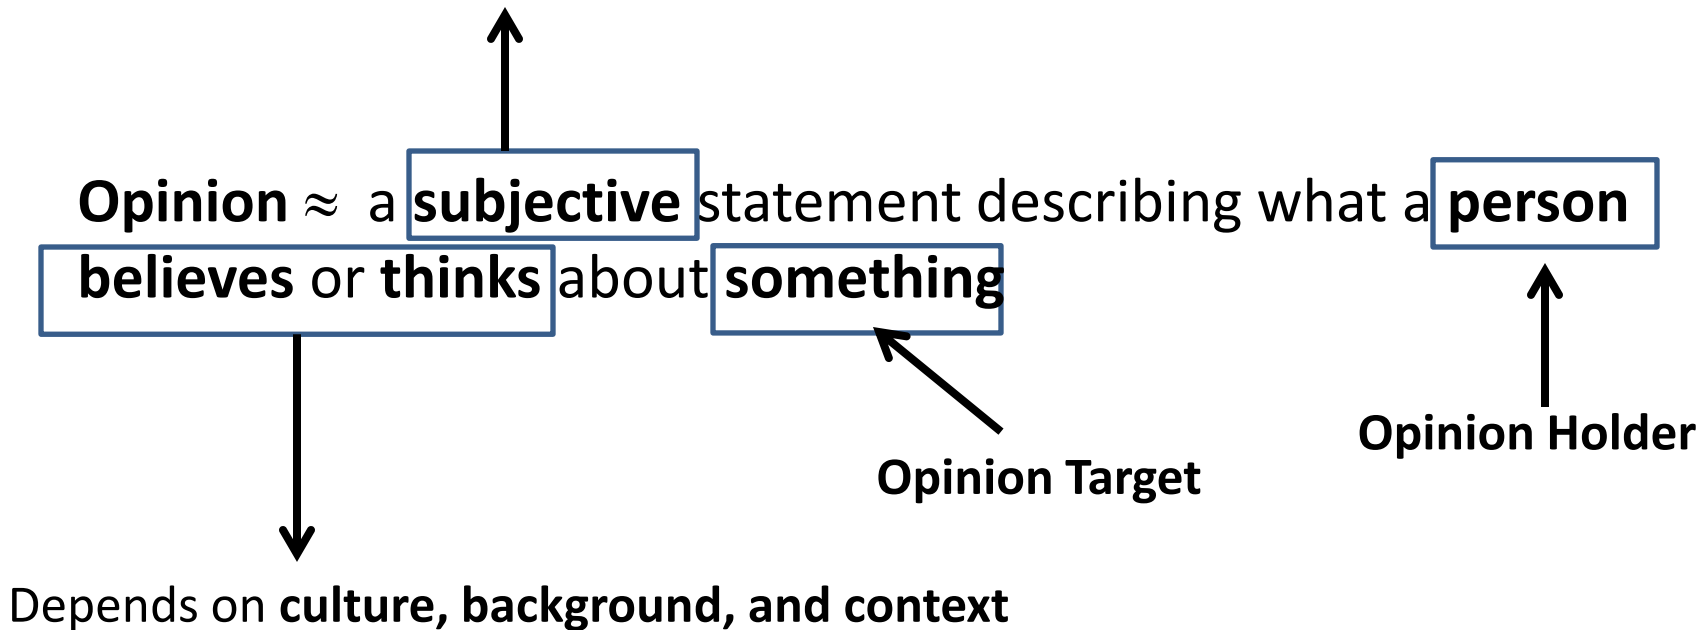
# Objective vs. Subjective Sensors



**How can we mine and analyze opinion buried in text?**

# What Is an Opinion?

**Objective** statement or **Factual** statement (can be proved right/wrong)



# Opinion Representation

- Basic Opinion Representation
  - Opinion **holder**: Whose opinion is this?
  - Opinion **target**: What is this opinion about?
  - Opinion **content**: What exactly is the opinion?
- Enriched Opinion Representation
  - Opinion **context**: Under what situation (e.g., time, location) was the opinion expressed?
  - Opinion **sentiment**: What does the opinion tell us about the opinion holder's feeling (e.g., positive vs. negative)?

# A Product Review (Explicit Holder and Target)

- Basic Opinion Representation

- Opinion **holder**: Whose opinion is this?

Reviewer X

- Opinion **target**: What is this opinion about?

Product: iPhone 6

- Opinion **content**: What exactly is the opinion?

Review Text

- Enriched Opinion Representation

- Opinion **context**: Under what situation (e.g., time, location) was the opinion expressed?

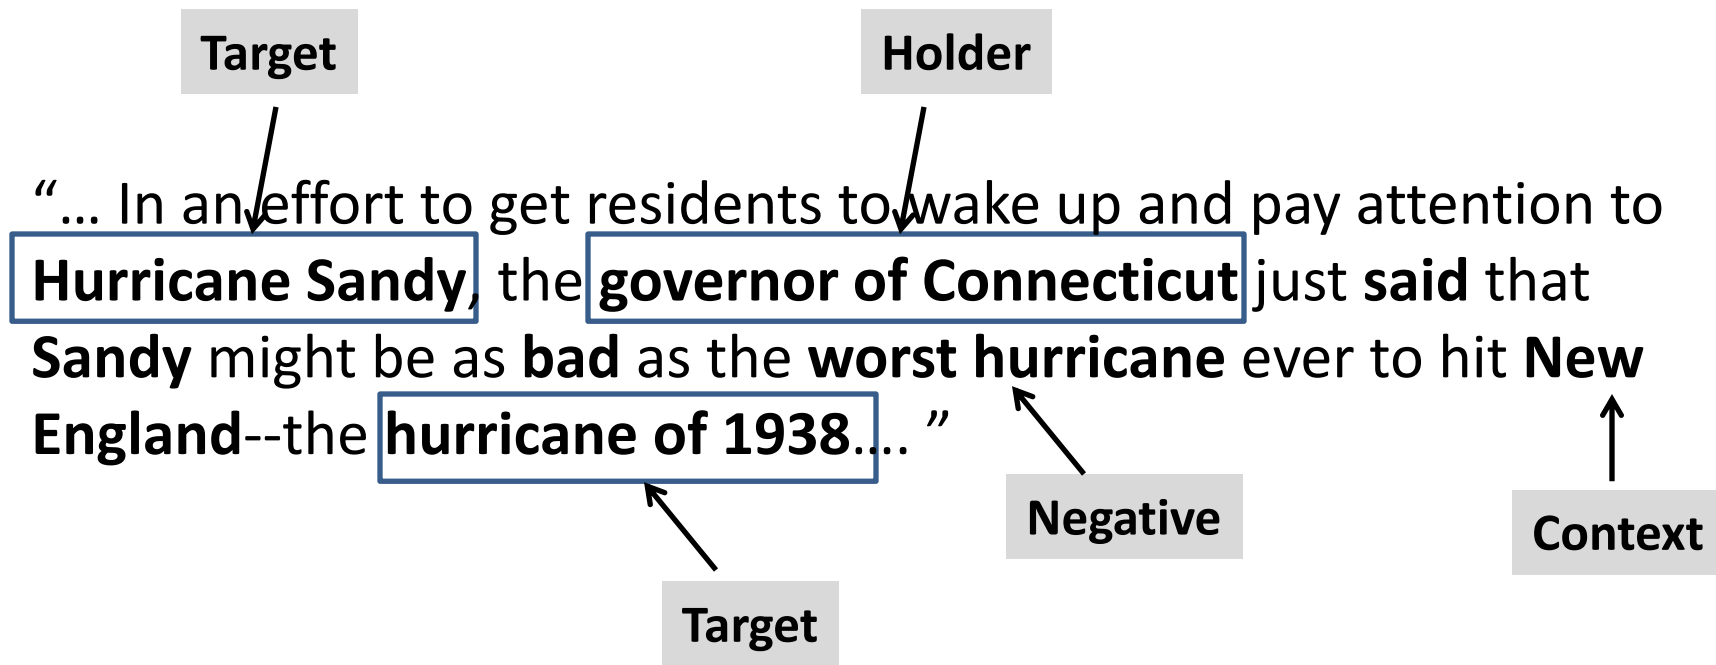
Year = 2015

- Opinion **sentiment**: What does the opinion tell us about the opinion holder's feeling (e.g., positive vs. negative)?

Positive

Relatively Easy to Mine and Analyze

# A Sentence in News (Implicit Holder and Target)



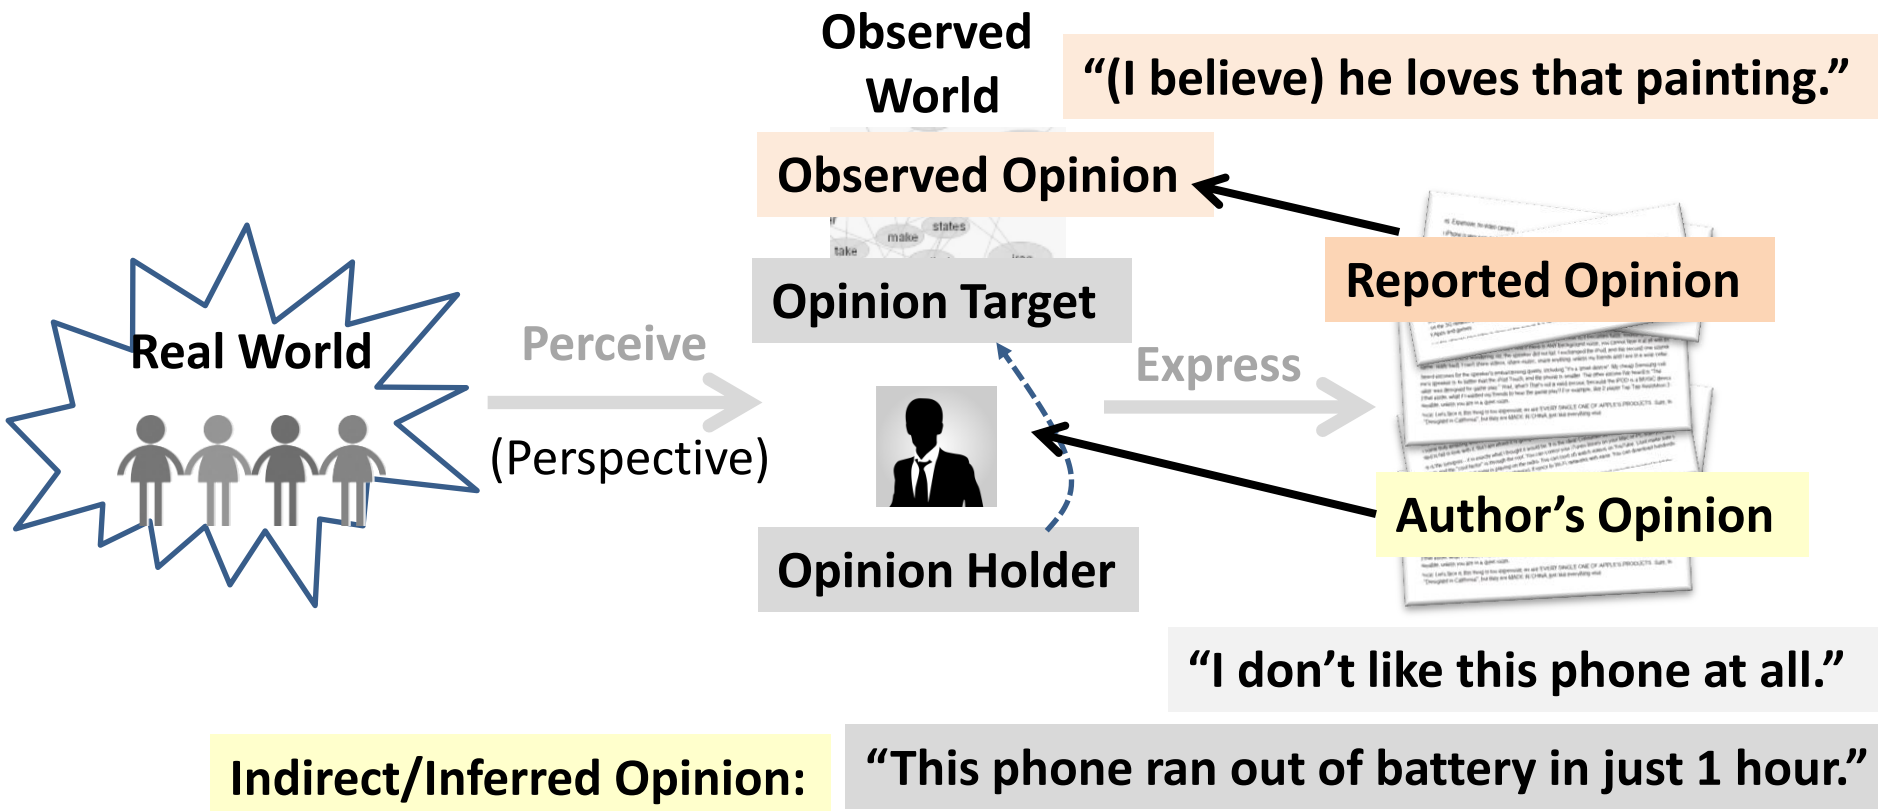
**Harder to Mine and Analyze: Need deeper NLP**



# Variations of Opinions

- **Opinion holder:** Individual vs. group
- **Opinion target:** One entity, a group of entities, one attribute of an entity, someone else's opinion, etc.
- **Opinion content:**
  - Surface variation: one sentence/phrase, a paragraph, a whole article
  - Sentiment/emotion variation: positive vs. negative, happy vs. sad, etc.
- **Opinion context**
  - Simple context: Different time, location, etc.
  - Complex context: Potentially includes the entire discourse context of an opinion

# Different Kinds of Opinions in Text Data



# The Task of Opinion Mining

## A Set of Opinion Representations

### Text Data



Opinion Holder

Opinion Target

Opinion Content

Opinion Context


Opinion  
Sentiment

Often some elements of the representation are already known

Simplest Opinion Mining task(s)?

# Why Opinion Mining?

- **Decision Support**
  - Help consumers choose a product or service
  - Help voters decide whom to vote for
  - Help policy makers design new policy
- **Understand People**
  - Help understand people's preferences to better serve them (e.g., optimize a product search engine; optimize recommender systems)
  - Help with advertising (targeted advertising)
- **“Voluntary Survey” (humans as sensors; aggregated opinions)**
  - Business intelligence
  - Market research
  - Data-driven social science research
  - Gain advantage in **any** prediction (text-based prediction)

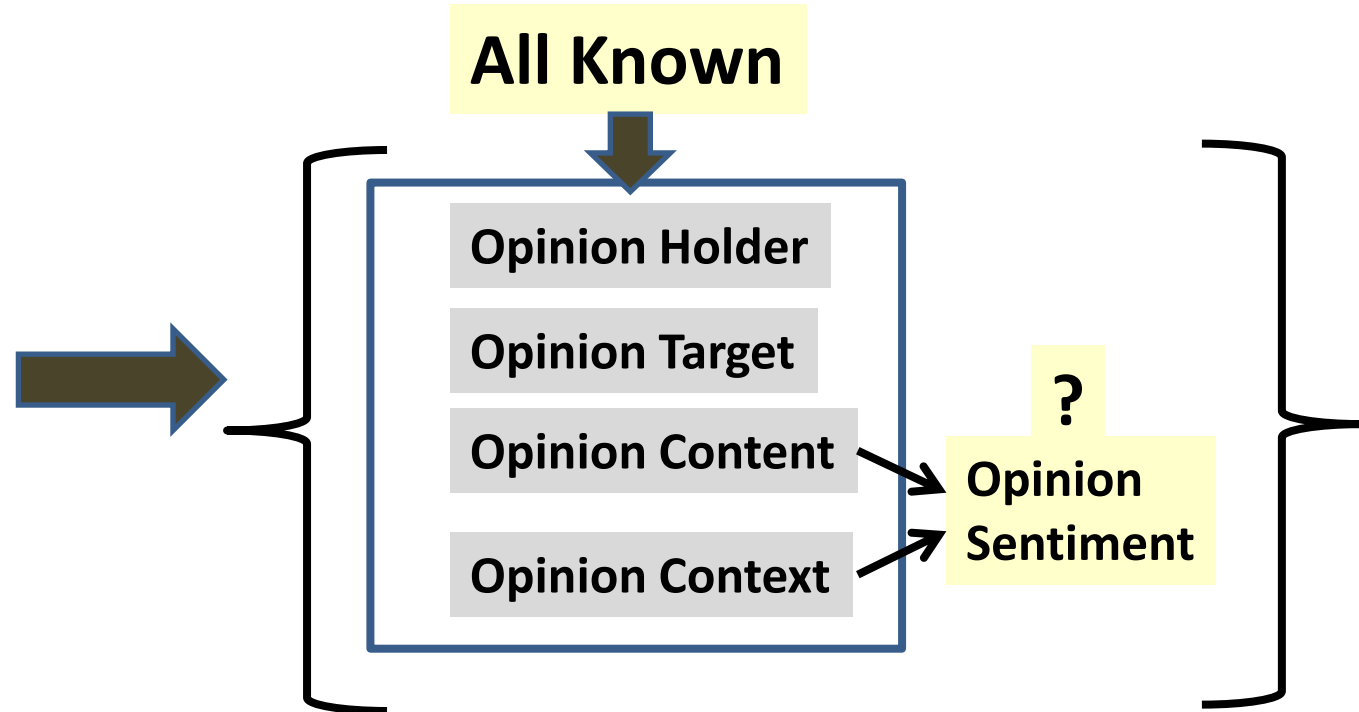


# Opinion Mining and Sentiment Analysis: Sentiment Classification

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# Sentiment Classification

# Text Data



# Sentiment Classification: Task Definition

- Input: An opinionated text object
- Output: A sentiment tag/label
  - Polarity analysis: e.g., categories = {positive, negative, neutral}, or categories = {5, 4, 3, 2, 1}
  - Emotion analysis (beyond polarity): e.g., categories = {happy, sad, fearful, angry, surprised, disgusted}
- A special case of text categorization! ➔ Any text categorization method can be used to do sentiment classification
- Further improvement comes from
  - More sophisticated features appropriate for sentiment tagging
  - Consideration of the order of the categories (e.g., ordinal regression)

# Commonly Used Text Features

- Character n-grams: can be mixed with different n's
  - General and robust to spelling/recognition errors, but less discriminative than words
- Word n-grams: can be mixed with different n's
  - Unigrams are often very effective, but not for sentiment analysis (e.g. , “it’s not good” or “it’s not as good as”)
  - Long n-grams are discriminative, but may cause overfitting
- POS tag n-grams: mixed n-gram with words and POS tags
  - E.g., “ADJECTIVE NOUN” or “great NOUN”




# Commonly Used Text Features (cont.)

- Word classes
  - Syntactic (= POS tags)
  - Semantic Concept: e.g., thesaurus/ontology, recognized entities
  - Empirical word clusters (e.g., cluster of paradigmatically or syntagmatically related words)
- Frequent patterns in text (e.g., frequent word set; collocations)
  - More specific/discriminative than words
  - May generalize better than pure n-grams
- Parse tree-based (e.g., frequent subtrees, paths)
  - Even more discriminative, but need to avoid overfitting
- Pattern discovery algorithms are very useful for feature construction

# NLP Enriches Text Representation with Complex Features

A dog is chasing a boy on the playground




**Dog(d1). Boy(b1). Playground(p1). Chasing(d1,b1,p1).**

**Speech Act = REQUEST**

“great NOUN”  
“Verb Adv Adj”

...



# Feature Construction for Text Categorization

- Feature design affects categorization accuracy significantly
- A combination of machine learning, error analysis, and domain knowledge is most effective
  - Domain knowledge → seed features, feature space
  - Machine learning → feature selection, feature learning
  - Error analysis → feature validation
- NLP enriches text representation → enriches feature space (more likely overfitting!)
- Optimizing the tradeoff between **exhaustivity** and **specificity** is a major goal

high coverage (frequent)

discriminative (infrequent)



# Sentiment Analysis: Ordinal Logistic Regression

ChengXiang “Cheng” Zhai  
Department of Computer Science  
University of Illinois at Urbana-Champaign

# Motivation: Rating Prediction

- Input: An opinionated text document  $\mathbf{d}$
- Output: Discrete rating  $\mathbf{r} \in \{1, 2, \dots, k\}$
- Using regular text categorization techniques
  - Doesn't consider the order and dependency of the categories
  - The features distinguishing  $r=2$  from  $r=1$  may be the same as those distinguishing  $r=k$  from  $r=k-1$  (e.g., positive words generally suggest a higher rating)
- Solution: Add order to a classifier (e.g., ordinal logistic regression )

# Logistic Regression for Binary Sentiment Classification

**Binary Response Variable:**  $Y \in \{0,1\}$       **Predictors:**  $X = (x_1, x_2, \dots, x_M)$ ,  $x_i \in \mathbb{R}$

$$Y = \begin{cases} 1 & X \text{ is POSITIVE} \\ 0 & X \text{ is NEGATIVE} \end{cases}$$

$$\log \frac{p(Y = 1 | X)}{p(Y = 0 | X)} = \log \frac{p(Y = 1 | X)}{1 - p(Y = 1 | X)} = \beta_0 + \sum_{i=1}^M x_i \beta_i \quad \beta_i \in \mathbb{R}$$


$$p(Y = 1 | X) = \frac{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i}}{e^{\beta_0 + \sum_{i=1}^M x_i \beta_i} + 1}$$

# Logistic Regression for Multi-Level Ratings

$$Y_j = \begin{cases} 1 & \text{rating is } j \text{ or above} \\ 0 & \text{rating is lower than } j \end{cases}$$

**Predictors:**  $X = (x_1, x_2, \dots, x_M)$ ,  $x_i \in \mathcal{R}$

**Rating:**  $r \in \{1, 2, \dots, k\}$



$$\log \frac{p(Y_j = 1 | X)}{p(Y_j = 0 | X)} = \log \frac{p(r \geq j | X)}{1 - p(r \geq j | X)} = \alpha_j + \sum_{i=1}^M x_i \beta_{ji} \quad \beta_{ji} \in \mathcal{R}$$

$$p(r \geq j | X) = \frac{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}}}{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}} + 1}$$

# Rating Prediction with Multiple Logistic Regression Classifiers

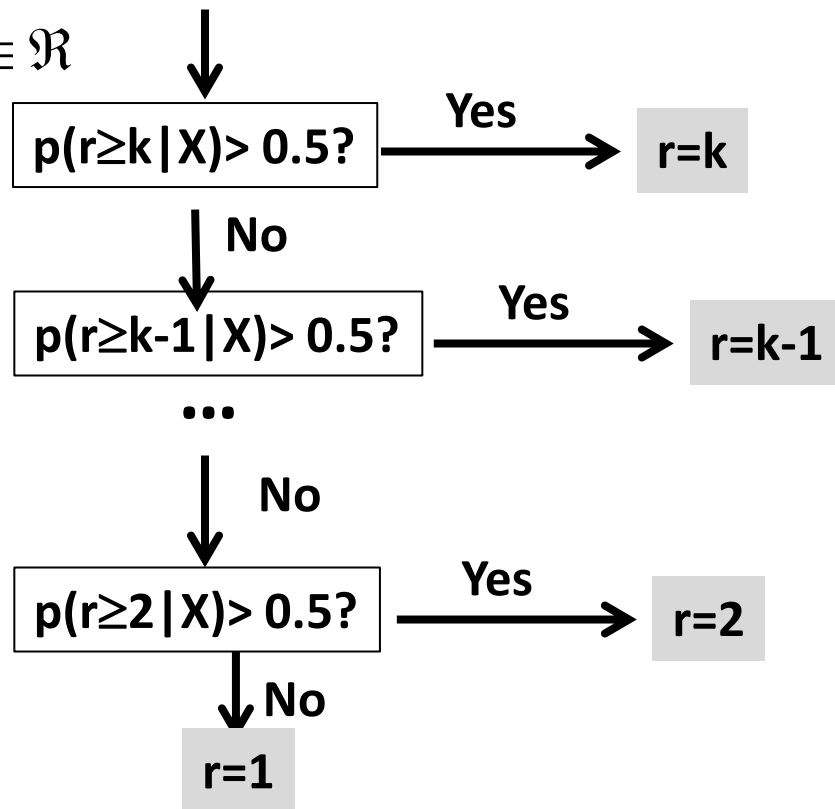
**Text Object:**  $X = (x_1, x_2, \dots, x_M)$ ,  $x_i \in \mathcal{R}$

**Rating:**  $r \in \{1, 2, \dots, k\}$

After training  $k-1$   
Logistic Regression Classifiers

$$p(r \geq j | X) = \frac{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}}}{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}} + 1}$$

$j = k, k-1, \dots, 2$






# Problems with k-1 Independent Classifiers?

$$\log \frac{p(Y_j = 1 | X)}{p(Y_j = 0 | X)} = \log \frac{p(r \geq j | X)}{1 - p(r \geq j | X)} = \alpha_j + \sum_{i=1}^M x_i \beta_{ji} \quad \beta_{ji} \in \Re$$

$$p(r \geq j | X) = \frac{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}}}{e^{\alpha_j + \sum_{i=1}^M x_i \beta_{ji}} + 1}$$



How many parameters are there in total?  **$(k-1) \cdot (M+1)$**

The k-1 classification problems are dependent.  
The positive/negative features tend to be similar!


# Ordinal Logistic Regression

**Key Idea:**  $\forall i = 1, \dots, M, \forall j = 3, \dots, k, \beta_{ji} = \beta_{j-1i}$

➔ Share training data

➔ Reduce # of parameters

$$\log \frac{p(Y_j = 1 | X)}{p(Y_j = 0 | X)} = \log \frac{p(r \geq j | X)}{1 - p(r \geq j | X)} = \alpha_j + \sum_{i=1}^M x_i \beta_i \quad \beta_i \in \mathcal{R}$$



$$p(r \geq j | X) = \frac{e^{\alpha_j + \sum_{i=1}^M x_i \beta_i}}{e^{\alpha_j + \sum_{i=1}^M x_i \beta_i} + 1}$$

How many parameters are there in total?

**M+k-1**

# Ordinal Logistic Regression: Rating Prediction

$$p(r \geq j | X) \geq 0.5 \Leftrightarrow \frac{e^{\alpha_j + \text{score}(X)}}{e^{\alpha_j + \text{score}(X)} + 1} \geq 0.5 \Leftrightarrow \text{score}(X) \geq -\alpha_j$$

