

Text Categorization: Evaluation

Part 1

ChengXiang “Cheng” Zhai
Department of Computer Science
University of Illinois at Urbana-Champaign

Overview

- What is text categorization?
- Why text categorization?
- How to do text categorization?
 - Generative probabilistic models
 - Discriminative approaches
- **How to evaluate categorization results?**

General Evaluation Methodology

- Have humans to create a test collection where every document is tagged with the desired categories (“ground truth”)
- Generate categorization results using a system on the test collection
- Compare the system categorization decisions with the human-made categorization decisions and quantify their similarity (or equivalently difference)
 - The higher the similarity is, the better the results are
 - Similarity can be measured from different perspectives to understand the quality of results in detail (e.g., which category performs better?)
 - In general, different categorization mistakes may have a different cost that inevitably depends on specific applications, but it is okay not to consider such a cost variation for **relative comparison of methods**

Classification Accuracy (Percentage of Correct Decisions)

	c₁	c₂	c₃	...	c_k	
d₁	y(+)	y(-)	n(+)		n(+)	+/- human answer
d₂	y(-)	n(+)	y(+)		n(+)	(+= correct; - =incorrect)
d₃	n(+)	n(+)	y(+)		n(+)	y/n system result
...						(y=yes; n=no)
d_N				

Classification Accuracy =

Total number of correct decisions

Total number of decisions made

=

count(y(+)) + count(n(-))

kN

Problems with Classification Accuracy

- Some decision errors are more serious than others
 - It may be more important to get the decisions right on some documents than others
 - It may be more important to get the decisions right on some categories than others
 - E.g., spam filtering: missing a legitimate email costs more than letting a spam go
- Problem with imbalanced test set
 - Skewed test set: 98% in category 1; 2% in category 2
 - Strong baseline: put all instances in category 1 → 98% accuracy!

Per-Document Evaluation

	c_1	c_2	c_3	\dots	c_k
d_1	y(+)	y(-)	n(+)		n(+)
d_2	y(-)	n(+)	y(+)		n(+)
d_3	n(+)	n(+)	y(+)		n(+)

How good are the decisions on d_i ?

When the system says “yes,”
how many are correct?

↓

Precision = $\frac{TP}{TP + FP}$

↗

Recall = $\frac{TP}{TP + FN}$

	System (“y”)	System (“n”)
Human (+)	True Positives TP	False Negatives FN
Human (-)	False Positives FP	True Negatives TN

Does the doc have all the categories
it should have?

Per-Category Evaluation

	c_1	c_2	c_3	...	c_k
d_1	y(+)	y(-)	n(+)		n(+)
d_2	y(-)	n(+)	y(+)		n(+)
d_3	n(+)	n(+)	y(+)		n(-)

How good are the decisions on c_i ?

When the system says “yes,”
how many are correct?

↓

Precision = $\frac{TP}{TP + FP}$

↗

Recall = $\frac{TP}{TP + FN}$

Has the category been assigned to
all the docs of this category?

	System (“y”)	System (“n”)
Human (+)	True Positives TP	False Negatives FN
Human (-)	False Positives FP	True Negatives TN

Combine Precision and Recall: F-Measure

$$F_{\beta} = \frac{1}{\frac{\beta^2}{\beta^2+1} \frac{1}{R} + \frac{1}{\beta^2+1} \frac{1}{P}} = \frac{(\beta^2 + 1)P * R}{\beta^2 P + R}$$

$$F_1 = \frac{2PR}{P + R}$$

P: precision

R: recall

β: parameter (often set to 1)

Why not $0.5 * P + 0.5 * R$?

What is R if the system says “y” for all category-doc pairs?